

# Lecture 1: Introduction

## Intelligent versus Automated

- **Automated:** run within a well-defined set of parameters and are very restricted in what tasks they can perform
- **Intelligent, autonomous systems:** self-governing, adapts to changes in the environment

Systems	Autonomous or Automated?
ATM	Automated
Disease outbreak detection	Intelligent
Kettle with automatic shut off	Automated
Self-driving cars	Intelligent
Warehouse robots	Insufficient Information

## Types of Intelligent Decisions

- **Single Step Decisions**
  - Deciding which is the "best" action to select for a given input
  - Output does not affect future input or output

- **Sequential Decision Making** Deciding which action to select for a given input or situation, considering future actions

# Lecture 2: Agents

**Agent:** anything that perceives its environment through sensors and acts on that environment through actuators

- PEAS Descriptions of Task Environments
- Performance, Environment, Actuators, Sensors

Example: Medical Diagnosis Agent	Task Environment	Observable	Deterministic	Episodic	Static	Discrete	Agents
Performance Measure	Environment	Actuators	Sensors				
Healthy patient, minimize costs, lawsuits	Patient, hospital, staff	Display questions, tests, diagnoses, treatments, referrals	Keyboard entry of symptoms, findings, patient's answers				

Properties of Environments	Rational Agent
<b>Fully observable:</b> can access complete state of environment at each point in time	<b>Stochastic:</b> when actions have multiple outcomes, each prescribed by a probability
<b>Deterministic:</b> next state of the environment completely determined by current state and agent's action	<b>Episodic:</b> agent's experience divided into independent, atomic episodes in which agent perceives and performs a single action in each episode
<b>Static:</b> agent doesn't need to keep sensing while decides what action to take, doesn't need to worry about time	<b>Dynamic:</b> environment changes while agent is thinking (changes with time)
<b>Discrete:</b> (note: applies to states, time, percepts, or actions)	<b>Continuous:</b> continuous values of states and/or actions
<b>Single agent:</b> single decision-making and executing entity	<b>Multiagent:</b> multiple decision-making/executing entities; cooperative or competitive

## Types of Agents

- **Simple Reflex Agent**
  - Selects actions using only the current percept
  - Works on condition-action rules: if condition then action
- **Model-based Reflex**
  - Maintain some internal state that keeps track of the part of the world it can't see now
  - Needs model

## Utility-directed Agents

- Utility measures which states are preferable to other states
- Assign numeric values to each possible outcome (utility or "happiness")
- Multidimensional utility (quality, failure rate, etc.)
- Time-dependent utility (hard/soft deadlines)
- Subjective vs. objective utility functions

## Learning Agents

- Successful agents split task of computing policy in 3 periods
- 1. Initially, designers compute some prior knowledge to include in policy
- 2. When deciding its next action, agent does some computation
- 3. Agent learns from experience to modify its behavior

Learn from experience to compensate for partial or incorrect prior knowledge

## Uninformed Search

- Breadth-first search (open list is FIFO queue)
- Depth-first search (open list is a LIFO queue)
- Uniform-cost search (shallowest node first)
- Depth-limited search (DFS with cutoff)
- Iterative-deepening search (incrementing cutoff)
- Bidirectional search (forward and backward)

## Lecture 3: Uninformed Search

### Control Types

Open loop control	Closed loop control
Decision depends on start and goal state	Decision depends on each state
No sensing at each state	Sensing at each state

### BFS

Expanded Nodes	Frontier List
(S)	(S)
(S)	(1,2)
(S,1,2)	(3,4,5,6)
(S,1,2,3)	(4,5,6)
(S,1,2,3,4)	(5,6)
(S,1,2,3,4,5)	(6)
(S,1,2,3,4,5,6)	()

### DFS

Expanded Nodes	Frontier List
(S)	(1,2)
(S,1)	(3,4,2)
(S,1,3)	(4,2)
(S,1,3,4)	(2)
(S,1,3,4,2)	(5,6)
(S,1,3,4,2,5)	(6)
(S,1,3,4,2,5,6)	()

# Lecture 4: Uninformed and Informed Search


## Depth-limited Search

Complete?	No (if shallowest goal node beyond depth limit)
Optimal?	No (if depth limit > depth of shallowest goal node and we expand a much longer path than the optimal one first)
Time Complexity	$O(b^l)$
Space Complexity	$O(b)$

- Solves infinite path problem by using predetermined depth limit  $l$
- Nodes at depth  $l$  are treated as if they have no successors
- Can use knowledge of the problem to determine  $l$  (but in general you don't know this in advance)

## Greedy Best-First Search

### Greedy Best-First Search Example



Corvallis -> McMinnville -> Wilsonville = 74 miles


### Evaluating Greedy Best-First Search

Complete?	Yes if the graph is finite and the heuristic function is informative (i.e. not 0 at all nodes).
Optimal?	No
Time Complexity	$O(b^m)$
Space Complexity	$O(b^m)$

Greedy Best-First search results in lots of unnecessary nodes being expanded

## Informed Search

### Greedy Best-First Search: Navigating in Manhattan



$g(n)$  = cost of path from the initial state to  $n$   
 $h(n)$  = estimate of the remaining distance

## Heuristic Search: A\*


$$f(n) = g(n) + h(n)$$

## Admissibility and Consistency

- **Admissible heuristic:** never overestimates the actual cost to reach a goal.

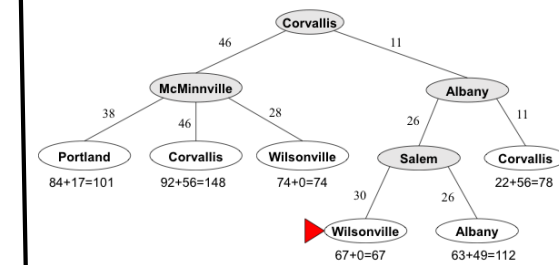
- **Consistent (or monotone) heuristic:**

$$h(n) \leq c(n, a, n') + h(n')$$



- **Every consistent heuristic is also admissible**


### A\* Search Example



Heuristic:  $h(n)$

Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

## Proper termination: Stop when you pop a goal state from the priority queue



# Local Search

1. Hill-climbing
2. Simulated Annealing
3. Beam Search
4. Genetic Algorithms
5. Gradient Descent

## Hill-climbing (Intuitively)

- Starting at initial state X, keep moving to the neighbor with the highest objective function value greater than X's.

### Hill Climbing Search

- Hill-climbing also called **greedy** local search
- Greedy because it takes the **best immediate move**
- Greedy algorithms often perform quite well
  - Disadvantage: all k states can become stuck in a small region of the state space
  - To fix this, use stochastic beam search
  - Stochastic beam search:
    - Doesn't pick best k successors
    - Chooses k successors at random, with probability of choosing a given successor being an increasing function of its value
- Problems:
  - Can get stuck at a **local maximum**.
  - Unable to find its way off a **plateau**.
  - Cannot climb along a narrow **ridge** when each possible step goes down.

## Simulated Annealing

- Hill-climbing never makes a downhill move
- What if we added some random moves to hill-climb help it get out of local maxima?
- This is the motivation for **simulated annealing**
- Generate successors randomly
- Allow "bad" moves with some probability
- How to select p?

### Genetic Algorithms

- Like natural selection in which an organism creates offspring according to its fitness for the environment
  - Essentially a variant of stochastic beam search that combines two parent states
  - Over time, population contains individuals with high fitness
- Local Beam Search Example**
- Travelling Salesperson Problem (k=2)
- Select the best k successors from the complete list and repeat the process
- An individual program is represented by a sequence of "genes".
  - The selection strategy is randomized with probability of selection proportional to "fitness".
  - Individuals selected for reproduction are randomly paired, certain genes are crossed-over, and some are mutated.

## Lecture 7: Adversarial Search

### The Minimax Value of a Node

The minimax value of a node is the utility for MAX of being in the corresponding state, *assuming that both players play optimally* from there to the end of the game

Minimax\_value(n) =

- UTILITY(n) If n is a terminal state
- $\max_{s \in \text{Successors}(n)} \text{Minimax\_value}(s)$  If n is a MAX node
- $\min_{s \in \text{Successors}(n)} \text{Minimax\_value}(s)$  If n is a MIN node

Minimax value maximizes worst-case outcome for MAX


### Properties

- Computes minimax decision from the current state
- Depth-first exploration of the game tree
- Complete?** Yes, if the graph is finite
- Optimal?** Yes, against an optimal opponent
- Time Complexity:**  $O(b^m)$  where b=# of legal moves, m=maximum depth of tree
- Space Complexity:**
  - $O(bm)$  if all successors generated at once
  - $O(m)$  if only one successor generated at a time (each partially expanded node remembers which successor to generate next)

### Alpha-Beta Pruning: Intuition

- If at a MIN player node, prune if minimax value of node  $\leq \alpha$
- If at a MAX player node, prune if minimax value of node  $\geq \beta$

iss Exercise #1: Alpha-Beta Pruning



## Lecture 8: Game Theory

### Dominant Strategies

- Normal form game: form of representing the game
- Strategies: different actions/ decision options
- Payoffs: utility of each decision
- Pure strategy: deterministic strategy selection
- Mixed strategy: probabilistic strategy selection

Suppose a player has two strategies S and S'. We say S **dominates** S' if choosing S always yields at least as good an outcome as choosing S'.

- S **strictly dominates** S' if choosing S always gives a better outcome than choosing S' (no matter what the other player does)
- S **weakly dominates** S' if there is one set of opponent's actions for which S is superior, and all other sets of opponent's actions give S and S' the same payoff.

### Example of Dominant Strategies

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

If Bob testifies, "testify" strongly dominates "refuse" strategy for Alice

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = 0, B = -1

If Bob refuses, "testify" weakly dominates "refuse" for Alice

Note

## Lecture 9: Game Theory

### Strategies and Equilibria

How to Spot a Nash Equilibrium

- Dominant strategy:** A player's best move, regardless of what other players do.
- Pareto optimality:** A state where no one can be made better off without making someone else worse off (i.e. the best for all players)
- Nash equilibrium:** A situation where no player can improve their outcome by changing their strategy, **while other players keep theirs the same**.
- Dominant strategy equilibrium:** A Nash equilibrium where all players have a dominant strategy.

	S1	S2	S3
S1	A = 0, B = 4	A = 4, B = 0	A = 5, B = 3
S2	A = 4, B = 0	A = 0, B = 4	A = 5, B = 3
S3	A = 3, B = 5	A = 3, B = 5	A = 6, B = 6

A won't change her Strategy of S3  
Payoff of 6 > 5 (S2) and 6 > 5 (S1)

B won't change his Strategy of S3  
Payoff of 6 > 5 (S2) and 6 > 5 (S1)

### Dominant Strategy Equilibrium

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

- (testify, testify) is a **dominant strategy equilibrium**
- It's an equilibrium because no player can benefit by switching strategies given that the other player sticks with the same strategy
- An equilibrium is a local optimum in the space of policies
- Pareto optimality:** A state where no one can be made better off without making someone else worse off (i.e. the best for all players)

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

	B: S1	B: S2	B: S3
A: S1	A=0, B=2	A=5, B=3	A=2, B=1
A: S2	A=1, B=3	A=2, B=1	A=7, B=4
A: S3	A=10, B=10	A=3, B=8	A=1, B=6

Does Player A have a strictly dominant strategy? If so, which one? **No**

	Best: cloud	Best: VR
ACME: cloud	A = 9, B = 9	A = -3, B = -1
ACME: VR	A = -4, B = -1	A = 5, B = 5

There are two Nash Equilibria in this game. In general, you can have multiple Nash Equilibria.

**Nash equilibrium:** A situation where no player can improve their outcome by changing their strategy, **while other players keep theirs the same**.

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

### Mixed Strategies

- Recall that a pure strategy is a deterministic policy i.e. you pick a strategy and play it all the time
- A **mixed strategy** is a randomized policy i.e. you select your strategy based on a probability distribution
- E.g. Select strategy S1 with probability p and strategy S2 with probability (1-p)
- Is there a mixed strategy Nash Equilibrium in 2 Fingered Morra?