

Lecture 1: Introduction

Intelligent versus Automated

- Automated: run within a well-defined set of parameters and are very restricted in what tasks they can perform
- Intelligent, autonomous systems: self-governing, adapts to changes in the environment

Systems	Autonomous or Automated?
ATM	Automated
Disease outbreak detection	Intelligent
Kettle with automatic shut off	Automated
Self-driving cars	Intelligent
Warehouse robots	Insufficient Information

Types of Intelligent Decisions

- Single Step Decisions
 - Deciding which is the "best" action to select for a given input
 - Output does not affect future input or output

- Sequential Decision Making
 - Deciding which action to select for a given input or situation, considering future actions

Lecture 2: Agents

Agent: anything that perceives its environment through sensors and acts on that environment through actuators

PEAS Descriptions of Task Environments

Performance, Environment, Actuators, Sensors

Example: Medical Diagnosis Agent

Performance Measure	Environment	Actuators	Sensors
Healthy patient, minimize costs, lawsuits	Patient, hospital, staff	Display questions, tests, diagnoses, treatments, patient's answers	Keyboard entry of symptoms, findings, patient's answers

Properties of Environments

Property	Description
Fully observable	can access complete state of environment at each point in time
Partially observable	could be due to noisy, inaccurate or incomplete sensor data
Deterministic	next state of the environment completely determined by current state and agent's action
Stochastic	when actions have multiple outcomes, each prescribed by a probability
Episodic	agent's experience divided into independent, atomic episodes in which agent perceives and performs a single action in each episode
Sequential	current decision affects all future decisions
Static	agent doesn't need to keep sensing while decides what action to take, doesn't need to worry about time
Dynamic	environment changes while agent is thinking (changes with time)
Discrete	(note: applies to states, time, percepts, or actions)
Continuous	continuous values of states and/or actions
Single agent	single decision-making and executing entity
Multiagent	multiple decision-making/executing entities; cooperative or competitive

Types of Agents

Simple Reflex Agent

Model-based Reflex

Utility-directed Agents

- Utility measures which states are preferable to other states
- Assign numerical values to each possible outcome (utility or "happiness")
- Multidimensional utility (quality, failure rate, etc.)
- Time-dependent utility (hard/soft deadlines)
- Subjective vs. objective utility functions

Learning Agents

- Successful agents split task of computing policy in 3 periods
- Initially, designers compute some prior knowledge to include in policy
- When deciding its next action, agent does some computation
- Agent learns from experience to modify its behavior

Learn from experience to compensate for partial or incorrect prior knowledge

Uninformed Search

- Breadth-first search (open list is FIFO queue)
- Depth-first search (open list is a LIFO queue)
- Uniform-cost search (shallowest node first)
- Depth-limited search (DFS with cutoff)
- Iterative-deepening search (incrementing cutoff)
- Bidirectional search (forward and backward)

Lecture 3: Uninformed Search

Control Types

Open loop control	Closed loop control
Decision depends on start and goal state	Decision depends on each state
No sensing at each state	Sensing at each state

BFS

Expanded Nodes	Frontier List
(S)	(S)
(S)	(1,2)
(S,1,2)	(3,4,5,6)
(S,1,2,3)	(4,5,6)
(S,1,2,3,4)	(5,6)
(S,1,2,3,4,5)	(6)
(S,1,2,3,4,5,6)	()

DFS

Expanded Nodes	Frontier List
(S)	(S)
(S)	(1,2)
(S,1)	(3,4,2)
(S,1,3)	(4,2)
(S,1,3,4)	(2)
(S,1,3,4,2)	(5,6)
(S,1,3,4,2,5)	(6)
(S,1,3,4,2,5,6)	()

Lecture 4: Uninformed and Informed Search

Depth-limited Search

Complete?	No (if shallowest goal node beyond depth limit)
Optimal?	No (if depth limit > depth of shallowest goal node and we expand a much longer path than the optimal one first)
Time Complexity	$O(b^l)$
Space Complexity	$O(b)$


- Solves infinite path problem by using predetermined depth limit

Nodes at depth l are treated as if they have no successors

Can use knowledge of the problem to determine l (but in general you don't know this in advance)

Greedy Best-First Search

Greedy Best-First Search Example



Corvallis -> McMinnville -> Wilsonville = 74 miles


Evaluating Greedy Best-First Search

Complete?	Yes if the graph is finite and the heuristic function is informative (i.e. not 0 at all nodes).
Optimal?	No
Time Complexity	$O(b^m)$
Space Complexity	$O(b^m)$

Greedy Best-First search results in lots of unnecessary nodes being expanded

Informed Search

Greedy Best-First Search: Navigating in Manhattan



$g(n)$ = cost of path from the initial state to n
 $h(n)$ = estimate of the remaining distance

Heuristic Search: A*


$$f(n) = g(n) + h(n)$$

Admissibility and Consistency

- Admissible heuristic: never overestimates the actual cost to reach a goal.

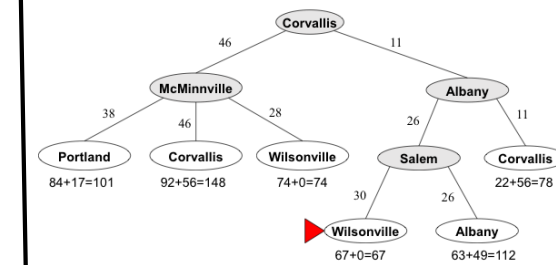
- Consistent (or monotone) heuristic:

$$h(n) \leq c(n, a, n') + h(n')$$



- Every consistent heuristic is also admissible


A* Search Example



Heuristic: $h(n)$

Corvallis	56
Albany	49
Salem	28
Portland	17
McMinnville	18

Proper termination: Stop when you pop a goal state from the priority queue



Summary of A* Search

Complete?	Yes if $h(n)$ is admissible, b is finite, and all step costs exceed some finite ϵ^1
Optimal?	Yes if $h(n)$ is admissible
Time Complexity	$O(b^d)$ (In the worst case but a good heuristic can reduce this significantly)
Space Complexity	$O(b^d)$

Local Search

1. Hill-climbing
2. Simulated Annealing
3. Beam Search
4. Genetic Algorithms
5. Gradient Descent

Hill-climbing (Intuitively)

- Starting at initial state X, keep moving to the neighbor with the highest objective function value greater than X's.

Hill Climbing Search

- Hill-climbing also called **greedy** local search
- Greedy because it takes the **best immediate move**
- Greedy algorithms often perform quite well

Problems:

- Can get stuck at a **local maximum**.
- Unable to find its way off a **plateau**.
- Cannot climb along a narrow **ridge** when each possible step goes down.

Simulated Annealing

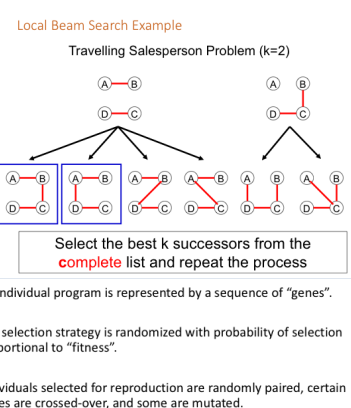
- Hill-climbing never makes a downhill move
- What if we added some random moves to hill-climb help it get out of local maxima?
- This is the motivation for **simulated annealing**
- Generate successors randomly
- Allow "bad" moves with some probability
- How to select p?

Genetic Algorithms

- Like natural selection in which an organism creates offspring according to its fitness for the environment
- Essentially a variant of stochastic beam search that combines two parent states
- Over time, population contains individuals with high fitness

6

Local Beam Search



Lecture 8: Game Theory

Dominant Strategies

- Normal form game: form of representing the game
- Strategies: different actions/ decision options
- Payoffs: utility of each decision
- Pure strategy: deterministic strategy selection
- Mixed strategy: probabilistic strategy selection

Suppose a player has two strategies S and S'. We say S **dominates** S' if choosing S always yields at least as good an outcome as choosing S'.

- S **strictly dominates** S' if choosing S always gives a better outcome than choosing S' (no matter what the other player does)
- S **weakly dominates** S' if there is one set of opponent's actions for which S is superior, and all other sets of opponent's actions give S and S' the same payoff.

Example of Dominant Strategies

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

If Bob testifies, "testify" strongly dominates "refuse" strategy for Alice

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = 0, B = -1

If Bob refuses, "testify" weakly dominates "refuse" for Alice

Note

Lecture 9: Game Theory

Strategies and Equilibria

How to Spot a Nash Equilibrium

- Dominant strategy:** A player's best move, regardless of what other players do.
- Pareto optimality:** A state where no one can be made better off without making someone else worse off (i.e. the best for all players)
- Nash equilibrium:** A situation where no player can improve their outcome by changing their strategy, while other players keep theirs the same.
- Dominant strategy equilibrium:** A Nash equilibrium where all players have a dominant strategy.

	S1	S2	S3
A	A = 0, B = 4	A = 4, B = 0	A = 5, B = 3
S2	A = 4, B = 0	A = 0, B = 4	A = 5, B = 3
S3	A = 3, B = 5	A = 3, B = 5	A = 6, B = 6

A won't change her Strategy of S3
Payoff of 6 > 5 (S2) and 6 > 5 (S1)

B won't change his Strategy of S3
Payoff of 6 > 5 (S2) and 6 > 5 (S1)

Lecture 7: Adversarial Search

The Minimax Value of a Node

The minimax value of a node is the utility for MAX of being in the corresponding state, assuming that both players play optimally from there to the end of the game

Minimax_value(n) =


- UTILITY(n) If n is a terminal state
- Max_{s ∈ Successors(n)} Minimax_value(s) If n is a MAX node
- Min_{s ∈ Successors(n)} Minimax_value(s) If n is a MIN node

Minimax value maximizes worst-case outcome for MAX

iss Exercise #1: Alpha-Beta Pruning

Properties

- Computes minimax decision from the current state
- Depth-first exploration of the game tree
- Complete?** Yes, if the graph is finite
- Optimal?** Yes, against an optimal opponent
- Time Complexity:** O(b^m) where b=# of legal moves, m=maximum depth of tree
- Space Complexity:**
 - O(bm) if all successors generated at once
 - O(m) if only one successor generated at a time (each partially expanded node remembers which successor to generate next)



Alpha-Beta Pruning: Intuition

function MAX-VALUE(state, α, β) returns a utility value
if TERMINAL-TEST(state) then return UTILITY(state)
v ← -∞
for each a in ACTIONS(state) do
v ← MAX(v, MIN-VALUE(RESULT(s,a), α, β))
if v ≥ β then return v
α ← MAX(α, v)
return v

function MIN-VALUE(state, α, β) returns a utility value
if TERMINAL-TEST(state) then return UTILITY(state)
v ← +∞
for each a in ACTIONS(state) do
v ← MIN(v, MAX-VALUE(RESULT(s,a), α, β))
if v ≤ α then return v
β ← MIN(β, v)
return v

max: $v \geq \beta$

min: $v \leq \alpha$

Dominant Strategy Equilibrium

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

- (testify, testify) is a dominant strategy equilibrium
- It's an equilibrium because no player can benefit by switching strategies given that the other player sticks with the same strategy
- An equilibrium is a local optimum in the space of policies
- Pareto optimality:** A state where no one can be made better off without making someone else worse off (i.e. the best for all players)

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

	B: S1	B: S2	B: S3
A: S1	A=0, B=2	A=5, B=3	A=2, B=1
A: S2	A=1, B=3	A=2, B=1	A=7, B=4
A: S3	A=10, B=10	A=3, B=8	A=1, B=6

Does Player A have a strictly dominant strategy? If so, which one? No

	Best: cloud	Best: VR
ACME: cloud	A = 9, B = 9	A = -3, B = -1
ACME: VR	A = -4, B = -1	A = 5, B = 5

There are two Nash Equilibria in this game. In general, you can have multiple Nash Equilibria.

Nash equilibrium: A situation where no player can improve their outcome by changing their strategy, while other players keep theirs the same.

	Bob: testify	Bob: refuse
Alice: testify	A = -5, B = -5	A = 0, B = -10
Alice: refuse	A = -10, B = 0	A = -1, B = -1

Mixed Strategies

- Recall that a pure strategy is a deterministic policy i.e. you pick a strategy and play it all the time
- A **mixed strategy** is a randomized policy i.e. you select your strategy based on a probability distribution
- E.g. Select strategy S1 with probability p and strategy S2 with probability (1-p)
- Is there a mixed strategy Nash Equilibrium in 2 Fingered Morra?